

ITU-MiniTwit

DevOps, Software Evolution and Software Maintenance
BSc Group F (gruppe 1)

Sebastian Holt Svendsen <sesv@itu.dk>
Andreas Nygaard Schiøtt <anys@itu.dk> Philip <xxxx@itu.dk>
Tobias <xxxx@itu.dk> Member Five <xxxx@itu.dk>

May 20, 2026

Word count:
2293 (errors:2)

Contents

1	System's Perspective	2
1.1	Architecture and Design	2
1.2	Dependencies	4
1.3	Current State of the Systems	5
2	Process' Perspective	7
2.1	CI/CD Pipeline	7
2.2	Monitoring	9
2.3	Logging	10
2.4	Security Hardening	11
2.5	Availability and Scaling	12
3	Reflection Perspective	13
3.1	Evolution and Refactoring	13
3.2	Maintenance	13
3.3	DevOps Style of Work	14
4	Use of Generative AI	15
4.1	Tools Used	15
4.2	Impact on Work Process	15
4.2.1	Supported	15
4.2.2	Hindered	15

[ITU-MiniTwit GitHub Repository](#)

1 System's Perspective

Authors: Sebastian Holt Svendsen

In this section, we provide an overview of the ITU-MiniTwit system from a technical perspective. We cover three areas: the architecture and design (section 1.1), the dependencies and technologies used (section 1.2), and the current state of the system (section 1.3). The app is deployed on DigitalOcean as a Docker Swarm cluster, fronted by two nginx load balancers, with a shared Redis for user-session-state and a managed PostgreSQL database for persistent data.

1.1 Architecture and Design

Figure 1 shows the production topology of ITU-MiniTwit deployed on DigitalOcean.




Figure 1: Deployment diagram — production topology of ITU-MiniTwit deployed on DigitalOcean.

We will be going through each tier of the diagram from top-to-bottom.

Tier 1: Public entry. Traffic enters via the DuckDNS domain, a free DNS service, resolving to a floating IP that can be reattached between the two load balancers.

Tier 2: Load balancers. Two nginx instances (LB1 MASTER, LB2 BACKUP) terminates TLS via Certbot and distribute requests to the next tier. Keepalived runs VRRP between them. On failure of the master, the backup promotes itself and the floating IP is reattached automatically.

Tier 3: App and data. The swarm contains two app nodes (App1, App2) each host one `itu-minitwit` replica plus a `node-exporter`. Swarm schedules them onto nodes labelled `role == app`. Both replicas share a managed PostgreSQL database and a Redis instance holding session state and ASP.NET DataProtection keys. Because the containers hold no per-request state, replicas are interchangeable and a failed node does not drop sessions.

Tier 4: Monitoring. A dedicated node runs Prometheus, Grafana, Loki, and Alloy. Prometheus scrapes `:9100` (host metrics) and `:5001` (app metrics). Alloy ships container logs from every node to Loki. Isolation ensures an app-tier outage does not blind the diagnostic tooling.

Application architecture. The application follows a layered design across three .NET projects with a one-way dependency rule. `Chirp.Core` (domain entities), `Chirp.Infrastructure` (EF Core `DbContext`, repositories, services, migrations), and `Chirp.Web` (Razor Pages scaffolded Identity pages, and the `MiniTwit-Stub` controller exposing the REST API).




Figure 2: Sequence diagram — request flow to the public timeline.

Figure 2 traces the runtime request flow. In the common case, the browser an `HTTPS GET`: nginx terminates TLS and forwards over `HTTP` to an app replica, which reads the caller's session from Redis, queries PostgreSQL, and returns the rendered page (`HTTP 200`). The `alt` fragment shows the alternative: with no valid session, the replica responds `HTTP 302 /login` and the browser follows the redirect. The diagram makes the responsibility split explicit, nginx only terminates TLS, the app is stateless, and Redis/PostgreSQL hold all shared state.

1.2 Dependencies

Tables 1 and 2 list the main technologies used in ITU-MiniTwit, grouped by the layer where they apply, from writing code through building, testing, deploying, and observing the system.

Table 1: Application dependencies, by layer.

Layer	Technology	Purpose
Language / runtime	C# / .NET 8 (ASP.NET Core)	Application language and web runtime
Web framework	ASP.NET Core Razor Pages	Server-rendered HTML UI
	ASP.NET Core Identity	Local accounts and authentication
	Swashbuckle / Swagger	Auto-generated docs for the simulator REST API
	prometheus-net.AspNetCore	Exposes the <code>/metrics</code> endpoint scraped by Prometheus
	Newtonsoft.Json	JSON serialisation for the simulator API
Data	Entity Framework Core 8 (Npgsql)	ORM; maps C# entities to tables and manages schema migrations
	PostgreSQL (DigitalOcean managed)	Persistent store for cheeps, authors and follow relations
	Redis (StackExchange.Redis)	Distributed session store and shared ASP.NET DataProtection key ring

Table 2: Infrastructure, build and operations dependencies, by layer.

Layer	Technology	Purpose
Containerisation	Docker Docker Swarm Docker Hub	Containerisation of all components Clusters multiple Docker hosts into one Registry for storing and distributing container images
Networking	nginx certbot keepalived DuckDNS	Terminates TLS and load balances to app replicas Automatically obtains and renews TLS certificates for the domain Implements VRRP for high availability of the load balancers Dynamic DNS service for the public hostname
Cloud infrastructure	DigitalOcean	Droplets for all cluster nodes, the managed PostgreSQL database, and the Floating IP
Provisioning / deployment	Shell scripts rsync over SSH	Bootstrapping, configuration, and deployment scripts for the cluster Copies the stack YAML files and monitoring configs to the swarm manager before deployment
CI/CD	GitHub Actions	Builds, tests and deploys the app on every push to main
Testing	xUnit Playwright	Unit and integration test framework for .NET End-to-end testing of the web UI
Static analysis	SonarCloud hadolint + ShellCheck	Code quality, coverage, and security analysis in CI Linting of Dockerfiles and shell scripts
Monitoring / logging	Prometheus Grafana Loki Grafana Alloy node-exporter	Metrics collection and storage Metrics and log visualisation Log aggregation and indexing Collects container logs and ships them to Loki Exposes host-level metrics for Prometheus
Dev tooling	Git / GitHub	Version control, code hosting, and pull-requests

1.3 Current State of the Systems

Static analysis. Code quality is tracked continuously through SonarCloud, which analyses every push to main. The current results are shown in fig. 9. Maintainability is rated **A** (54 minor issues) and duplication is low at **4.7%**, but reliability and security both sit at **C** (20 reliability issues, 5 security issues, and 22 security hotspots awaiting review). The **test coverage is only 13.1%**, the codebase’s clearest quality weakness. Dockerfiles, shell scripts and C# formatting are additionally linted in CI with hadolint, ShellCheck and `dotnet format`, configured to report findings without blocking the pipeline.




Figure 3: SonarCloud analysis summary for ITU-MiniTwit.

What works. In production the system is stable: the application reports healthy, serves the simulator’s load of around **2.2k req/s** with no in-flight backlog, uses about **156 MB** of memory, and returns its API endpoints in roughly **7 ms at the median and 37 ms at the 99th percentile** with no server (5xx) errors. High availability has been demonstrated end-to-end: keepalived failover between the two load balancers, two app replicas, automatic restart of crashed containers, and rolling updates without downtime and without logging users out. The full CI/CD pipeline (build, unit and integration tests, end-to-end tests, static analysis, and deployment) is in place, and the metrics and logging stacks are deployed. Over the term the project also migrated from SQLite to a managed PostgreSQL database and from per-node Docker Compose to Docker Swarm.

Open items. A number of limitations remain, documented here:

- The public-timeline page averages about **580 ms**, far slower than the API endpoints. It’s the main performance weakness, rooted in the database access pattern and round-trips to the managed database.
- The logging dashboard is dominated by Loki’s own internal push errors, so application logs are not effectively surfaced, and there is no alerting.
- SonarCloud’s security findings are unaddressed. The Docker images are not security-hardened and CI has no dedicated security scanner (e.g. CodeQL or Trivy).
- There is no declarative infrastructure-as-code or configuration management (Vagrant, Ansible). Provisioning is handled by shell scripts, and the **Vagrantfile** is a non-functional stub.
- With only two of five group members actively contributing, some weekly tasks (mapping the Three Ways, security scanning, infrastructure-as-code) were not reached.

2 Process' Perspective

Authors: Andreas Schjøtt

2.1 CI/CD Pipeline

Table 3: CI/CD tools

Technology	Purpose
GitHub Actions	CI/CD automation
Playwright	end-to-end testing
SonarCloud	Static analysis and code coverage
Hadolint	Dockerfile linter
ShellCheck	Shell script linter
Docker Build	Build and push container images to Docker Hub
Docker Swarm	Automated docker container orchestration
rsync	Idempotent Sync files to the remote swarm manager
Grafana	Logging / monitoring running Prometheus, loki, Alloy
Prometheus	Metric collection and monitoring
Loki	Log aggregation system
Grafana Alloy	Log and monitoring collection

Before cloud hosting, CI only built a docker container and ran unit tests against a local SQLite application.

After cloud hosting on **DigitalOcean** was setup we added the CD pipeline to the CI pipeline. The CI/CD pipeline were both triggered concurrently on push to main branch which then builds a **Docker image**, pushes it to **Docker Hub** and then deploys it to the DigitalOcean droplet using **SSH & docker pull**. With the later focus on idempotent deployment, we added a check to see if the CI pipeline succeeded before the CD pipeline is triggered.

The CI pipeline was extended to include end-to-end testing using **Playwright**, and static code analysis using **SonarCloud & Trivy**, **Hadolint** and **ShellCheck**. As shown in fig. 4, on every push the pipeline detects which paths changed, builds the project, and then runs static analysis tests, SonarCloud analysis, and end-to-end tests in parallel.

While writing the report we iteratively extended the CI pipeline to check which files were changed to decide which steps to run, if only the report directory was changed we only run "build PDF" report" step and vice versa. The pipeline also publishes a self-contained GitHub release each week.




Figure 4: CI pipeline (CI.yml).

The CD pipeline is triggered only after a successful CI run, ensuring that broken builds are never deployed. As shown in fig. 5, it consists of three sequential jobs: building and pushing the Docker image (*build-push*), deploying the application (*Deploy*), and deploying the monitoring stack (*Deploy-Monitoring*) with **Grafana**, **Prometheus**, **Loki** and **Grafana Alloy**.




Figure 5: CD pipeline

The *Deploy* job applies the Docker Swarm stack file via `docker stack deploy`. Rolling-update semantics are covered in section 1.1. Detailed steps in fig. 6




Figure 6: Detailed steps of the *Deploy* job.

2.2 Monitoring

[Link to Grafana monitoring dashboard.](#) Building on the monitoring topology in section 1.1, Prometheus scrapes the application's `metrics` endpoint (job `itu-minitwit-app`) and Node exporter on each Swarm node. Grafana renders the dashboards.

We have iteratively updated which data is collected and visualised in the Grafana dashboard based on the needs of the project. At the project's end we have the Grafana dashboard organised into five sections, as shown in fig. 7:

- **Health overview** — application up/down status, average response time, requests / sec, in-flight requests, error rate and current memory usage.
- **HTTP traffic** — request rate by status code, response-time percentiles grouped by p50 p90 and p99. request rate /s grouped by HTTP method, and the response time grouped by api endpoints.
- **Error tracking** error rates req/s grouped by 400/500 http response, errors measured in req/s per endpoint.
- **Process & runtime** — process CPU usage, process memory usage, thread count, garbage collection rate.
- **Uptime & Availability** — uptime & availability over time (bar graf), Availability measured in percentage over the selected time period.



Figure 7: Grafana monitoring dashboard.

2.3 Logging

[Link to Grafana logging dashboard.](#) Grafana Alloy runs on every Swarm node, reads the docker socket, and ships container logs to Loki labelled with container, service, stream, and image. Loki rejects samples older than 168 hours (7 days) to bound storage.

The application logs the following events:

- HTTP requests — method, path, response status code, and elapsed time measured in ms.
- Unhandled exceptions — method, path, error details.
- Custom logs through **ILogger** throughout the minitwit codebase.

The logs are visualised in a Grafana dashboard with the following panels:

- **Total Log lines & error lines** — Total log lines & errors within the selected time period
- **Log lines & Errors/exceptions & warnings per container** — Log lines & errors grouped by container within the selected time period
- **log volume over time** — Log lines grouped by time (bar graph) within the selected time period
- **log & Error viewer** — Visulisation of all log lines and errors.




Figure 8: Loki log aggregation in Grafana.

2.4 Security Hardening

We did not have the opportunity to implement hardened docker images as it is a paid feature in Docker Hub, we instead tried to implement parts of the guide [OWASP Docker Security Cheat Sheet](#) measures applied from Cheatsheet:

- **Unprivileged users in Docker containers** — Containers run as an unprivileged user to prevent privilege escalation attacks.
- **Prevent in-container privilege escalation** — use `--security-opt=no-new-privileges` on containers to prevent privilege escalation attacks.
- **Automatic Docker container update** - implementet a WeeklyHostUpdate which automatically updates the docker images on the swarm nodes to ensure that security updates are applied.

System uses applied to harden the system:

- **GitHub Secrets** — database connection strings and API keys are stored as encrypted secrets and injected at deploy time.
- **Firewall** — inbound and outbound traffic on both application droplets and the database is restricted to only the ports and sources required.
- **TLS/HTTPS** — traffic to the application is encrypted using TLS Through the Nginx load balancer using Certbot.
- **Static code analysis** — SonarCloud and trivy are used to identify security vulnerabilities in the codebase/pipelines/docker images.

Static code analysis is integrated in the CI pipeline using SonarCloud & Trivy. Every build is analysed for bugs, coverage, code duplication and security vulnerabilities. As shown in fig. 9, there are still a number of security vulnerabilities but with only two of five members actively contributing, there was not time to address them before the deadline




Figure 9: SonarCloud quality gate results.

2.5 Availability and Scaling

Authors: Andreas Schjøtt, Sebastian Holt Svendsen

The application topology described in section 1.1 already delivers our availability guarantees: keepalived VRRP between the two load balancers, two app replicas behind `least_conn` routing, Swarm's `on-failure` restart policy for crashed containers, and rolling updates with a 10-seconds delay between replicas to keep at least one instance serving traffic.

3 Reflection Perspective

Authors: Andreas Schjøtt, Sebastian Holt Svendsen

3.1 Evolution and Refactoring

Our initial implementation of CI/CD pipelines ran concurrently: a push would trigger deployment before its CI run had finished, so a failing build could reach production silently. We fixed this by gating CD on a successful CI run, and later added caching to speed up the End-to-End tests, which had been a bottleneck.

Early on we did not cull old containers on deployment. The resulting CPU pressure was first surfaced through our monitoring stack, which then became a crucial diagnostic tool. Monitoring and logging initially ran on the same droplet as the application, a single point of failure so we moved them onto a dedicated droplet, a decision that proved fortunate when that droplet was later compromised.

Hacking incident. On 2026-04-21 we discovered that our monitoring droplet had been compromised three days earlier: a brute-force SSH scanning campaign, ongoing since 2026-04-19, had culminated in the `preprod-deploy` key being used for root access from four Azure IPs. The attacker deployed a cryptominer disguised as `/tmp/mysql`, created a fake `dnsmasq` user for persistence, and started a rogue `chirp-postgres` container. Ironically the compromise was discovered not through our monitoring but because of it, the droplet was unresponsive and 100% CPU usage, prompting manual investigation. With root access via the leaked key, the host could not be salvaged. We destroyed and rebuilt it, restricting SSH access via UFW to our own IPs. Full timeline in [security-incident-ie-BITCOIN-MINE.md](#).

Architectural refactor. The system underwent several major architectural refactors over the semester, from a locally hosted single-container setup, to a Dockerised application using SQLite, to a DigitalOcean managed PostgreSQL, and finally to a Docker Swarm stack with replicated app instances behind a load-balanced floating IP. Each step was forced by a hard limit of the previous one: SQLite's concurrent write limit, single host failure modes that took the whole system down, and the need for rolling updates without downtime.

Across these refactors operational visibility was the common thread. Every significant improvement to the system responded to a bottleneck or failure surfaced either manually or through monitoring.

3.2 Maintenance

We did not implement an alerting system, which left a critical observability gap: the production environment could fail without anyone knowing until we manually checked. The consequences of this were most evident in the security incident described in section 3.1. Implementing an alerting system (at minimum Grafana alerts on uptime, CPU saturation, and authentication failures) would be the highest priority if we continued operating the system. Another instructive incident was self-inflicted. To resolve auth cookies being silently invalidated when load-balanced across replicas, we added `AddDataProtection().PersistKeysToStackExchangeRedis(...)` in `Program.cs`. The fix introduced a new failure mode: `ConnectionMultiplexer.Connect()` is blocking and eager, so the entire application failed to start if Redis was momentarily unavailable. This broke production on first deploy and our E2E test job (which had no Redis instance). We resolved it by adding `abortConnect=false` to the connection string and adding Redis to the E2E workflow ([commit46e5205](#)).

Lesson being fixes depending on additional infrastructure must be tested against that infrastructure's absence. Both incidents share the same shape: failures detectable only after the fact,

because we built observability for diagnosis but not for prevention. Closing the alerting gap and tightening test coverage of new infrastructure dependencies would be the highest maintenance priorities.

3.3 DevOps Style of Work

Compared to previous software projects, our development process was substantially more iterative and incremental, with continuous integration and delivery as the default rather than an afterthought. The most consequential shift was treating observability as a first-class feature. We built monitoring and logging early, and tried using them as our primary feedback loop. This habit has carried over into our other projects and courses.

A DevOps team is most effective when ownership, on-call responsibility, and review duties are distributed across all members. Ours was not. Of our five-person group, two of us carried the substantive work for the full semester. The others didn't commit much code, take branches, or return to tasks taken on. This was an unfortunate outcome of the group dynamics. Some of this we own: in the first weeks we moved quickly and likely raised the bar for participation higher than what was inclusive, without doing the team-formation work (shared expectations, paired tasks, defined responsibilities, etc.) that would have enabled on-ramping.

The operational cost was real: a bus-factor of two on every piece of infrastructure, the missing alerting in section 3.2 partly because no one else had the context to build it, and concentrated context that made routine handover or rotation impossible.

4 Use of Generative AI

Authors: Andreas Schjøtt, Sebastian Holt Svendsen

4.1 Tools Used

- Claude / Claude Code(CLI) For Brainstorming, debugging, Understanding course material, code generation, documentation, code reviews/Refactoring, process improvement suggestions.
- ChatGPT (OpenAI) For Brainstorming, debugging, Understanding course material, code generation, documentation, code reviews/Refactoring, process improvement suggestions.
- GitHub Copilot for code suggestions and autocompletion

4.2 Impact on Work Process

4.2.1 Supported

It helped us understand course material and concepts better by providing explanations and examples / deepening the explanations from the lectures/course material. able to get less stuck on problems, and get suggestions for how to solve them, which sped up development. helped with code generation for ci/cd pipelines, dockerfiles, and other boilerplate code.

4.2.2 Hindered

As mentioned earlier, we have been 2 people working on the project, this resulted in us being behind on the project, which meant that the use of generative AI could sometimes could become a crutch to catch up. To get value of the use of generative AI had to give alot of context to get usefull suggestions / answers /avoid hallucinations which made output useless. This has been helped by the use of CLI based tools like Claude Code/ Codex which could take the context from the codebase.